



Contents lists available at ScienceDirect



Big Data Research

www.elsevier.com/locate/bdr

SMR: Medical Knowledge Graph Embedding for Safe Medicine Recommendation

Fan Gong^a, Meng Wang^{b,c,*}, Haofen Wang^d, Sen Wang^e, Mengyue Liu^f^a Shanghai Shuguang Hospital Affiliated to Shanghai University of Traditional Chinese Medicine, Pu'an Road, Shanghai, China^b School of Computer Science and Engineering, Southeast University, Nanjing, China^c Key Laboratory of Computer Network and Information Integration (Southeast University), Ministry of Education, Nanjing, China^d College of Design and Innovation, Tongji University, Shanghai, China^e The University of Queensland, Brisbane, Australia^f School of Electronic and Information Engineering, Xi'an Jiaotong University, Xi'an, China

ARTICLE INFO

Article history:

Received 14 June 2020

Received in revised form 12 November 2020

Accepted 29 November 2020

Available online xxxx

Keywords:

Knowledge graph

Embeddings

Recommendation system

Drug safety

ABSTRACT

Most of the existing medicine recommendation systems that are mainly based on electronic medical records (EMRs) are significantly assisting doctors to make better clinical decisions benefiting both patients and caregivers. Even though the growth of EMRs is at a lighting fast speed in the era of big data, content limitations in EMRs restrain the existed recommendation systems to reflect relevant medical facts, such as drug-drug interactions. Many medical knowledge graphs that contain drug-related information, such as DrugBank, may give hope for the recommendation systems. However, the direct use of these knowledge graphs in systems suffers from robustness caused by the incompleteness of the graphs. To address these challenges, we stand on recent advances in graph embedding learning techniques and propose a novel framework, called Safe Medicine Recommendation (SMR), in this paper. Specifically, SMR first constructs a high-quality heterogeneous graph by bridging EMRs (MIMIC-III) and medical knowledge graphs (ICD-9 ontology and DrugBank). Then, SMR jointly embeds diseases, medicines, patients, and their corresponding relations into a shared lower dimensional space. Finally, SMR uses the embeddings to decompose the medicine recommendation into a link prediction process while considering the patient's diagnoses and adverse drug reactions. Extensive experiments on real datasets are conducted to evaluate the effectiveness of proposed framework.

© 2020 Published by Elsevier Inc.

1. Introduction

Over the last few years, medicine recommendation systems have been developed to assist doctors in making accurate medicine prescriptions. On the one hand, many researchers [1,2] adopt rule-based protocols that are defined by the clinical guidelines and the experienced doctors. Constructing, curating, and maintaining these protocols are time-consuming and labor-intensive. Rule-based protocols might be effective for a general medicine recommendation for a specific diagnosis, but give little help to tailored recommendations for complicated patients. On the other hand, supervised learning algorithms and variations, such as Multi-Instance Multi-label (MIML) learning [3], have been proposed to recommend medicines for patients. Both input features and ground-truth information that are extracted from massive EMRs are trained to obtain

* Corresponding author at: School of Computer Science and Engineering, Southeast University, Nanjing, China.

E-mail address: meng.wang@seu.edu.cn (M. Wang).

<https://doi.org/10.1016/j.bdr.2020.100174>

2214-5796/© 2020 Published by Elsevier Inc.

1
2
3
4
5
6
7
8
9
10
11
12

Fig. 1. Left part is a medical knowledge graph with missing relationships. Right part is prescription records in EMRs, and each edge indicates that a patient takes a medicine. There is no relationship between medicines in EMRs.

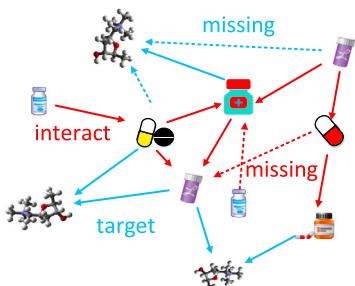
possible reason might because there is little medical expert knowledge in EMRs. Content limitations in EMRs constrain the systems to barely associate accurate medical facts with the recommended prescriptions, which makes the final recommendation less trustworthy for the complicated patients.

With the increasing emergence of knowledge graphs, many world-leading researchers have successfully extracted information from huge volumes of medical databases to build up giant heterogeneous graphs that reflect medical facts of medicines and diseases. For instance, DrugBank [8] is a rich source of medicine information. It contains extensive entities (drugs, drug targets, chemistry, etc.) and relationships (enzymatic pathways, drug-drug interactions, etc.). ICD-9 ontology [9] represents a knowledge base of human diseases and can be used to classify diagnoses of patients. Harnessing well-built medical knowledge graphs in EMRs-based medicine recommendation system might enable the reinvented system to provide appropriate prescriptions for special patients, as well as alerts of possible side effects and serious drug-drug interactions (DDIs).

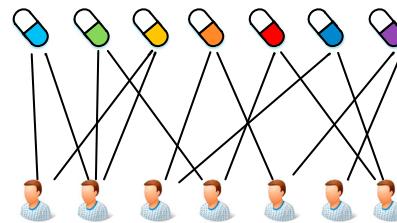
As shown in Fig. 1, linking EMRs and medical knowledge graphs to generate a large and high-quality heterogeneous graph is a promising pathway for medicine recommendations in a wider scope, but never easy. Specifically, the newly designed system confronts with the following challenges: **1. Computational Efficiency.** Querying specialized medical entities and relationships based on conventional graph-based algorithms have limitations in portability and scalability. The computational complexity becomes unfeasible when the heterogeneous graph reaches a very large scale. **2. Data Incompleteness.** The medical knowledge graphs also follow the long-tail distribution as same as other types of large-scale knowledge bases. Data incompleteness is another serious problem existing among entities and relationships in such a distribution. For example, since the DDIs is not usually identified in the clinical trial phase, there is a lack of significant DDIs in DrugBank which cannot support a comprehensive precaution to the medication. Last but not least, medicine recommendation suffers from **3. Cold Start.** As conventional systems normally recommend medicines based on the historical records, the pace of recommendation changes cannot keep up with the frequent updates of new therapies and treatments in medical practices. Little information on adverse reactions to the newly updated medicines in historical EMRs or even in well-built knowledge graphs makes the evidence-based recommendation model hardly support the new medicines as updated recommendations.

Taking all the challenges above into account, we propose a novel medicine recommendation framework based on graph embedding techniques, inspired by the idea of link prediction. We name our framework as Safe Medicine Recommendation (SMR) throughout this paper. The recommendation process mainly includes:

Medical Knowledge Graph



Electronic Medical Records



1. A large heterogeneous graph is constructed from EMRs and medical knowledge graphs, where the nodes are entities (medicines, diseases, patients), and the edges (or links) represent various relations between entities, such as drug-drug interactions.
2. The different parts of the generated heterogeneous graph (patient-medicine bipartite graph, patient-disease bipartite graph, medicine knowledge graph, disease knowledge graph) are embedded into a shared low-dimension space based on graph-based embedding models. Afterward, a joint learning algorithm is proposed to optimize the integrated graph simultaneously.
3. Based on the learned embeddings, a new patient, represented by the vectors of his/her diagnoses, is modeled as an entity in the disease-patient graph. Recommending medicines for the patient is translated to predict links from the patient to medicines.

The primary contributions of this work are summarized as follows:

1. We have developed graph-based embedding models to learn the effective representations of patients, diseases, and medicines in a shared low-dimension space. The representation of medicines enables the proposed framework can even effectively recommend newly emerged medicines for patients, which distinguishes most of the existing works.
2. To recommend safe medicines for new patients, we propose a novel method for modeling a patient based on the learned graph embeddings and make a safe recommendation by minimizing the potential adverse drug reactions.
3. We have conducted extensive experiments on large real-world datasets (MIMIC-III, DrugBank, and ICD-9 ontology) to evaluate the effectiveness of our framework. The experimental results have shown that the proposed framework outperforms all the compared methods.
4. To our best knowledge, we firstly propose a framework to conduct Safe Medicine Recommendation (SMR) and formulate it as a link prediction problem. The implementation generates a high-quality heterogeneous graph in which relationships among patients, diseases, and medicines can be unveiled in a wider scope.

The remainder of this paper is organized as follows: Section 2 details our proposed framework SMR. Section 3 reports the experimental results and Section 4 reviews related work. Section 5 presents the conclusions and future work.

67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132

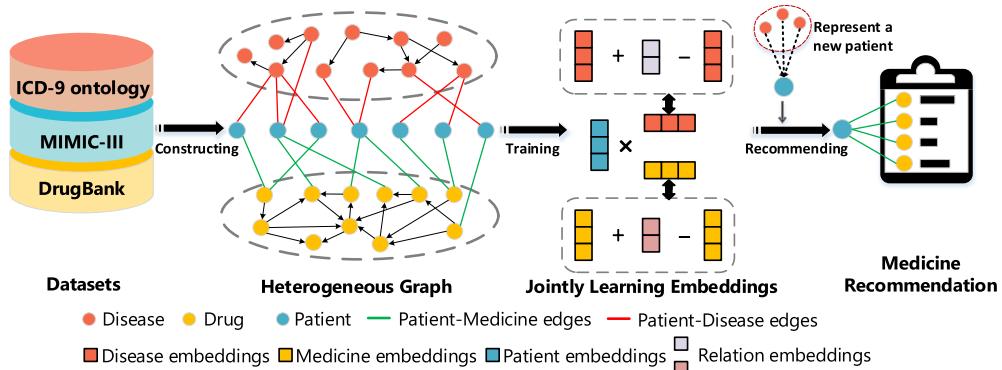


Fig. 2. Overview of our framework. Patient-disease and patient-medicine graphs are all bipartite graphs, while disease and medicine graphs are general graphs. Patient, disease, and medicine are encoded into a low dimensional metric by graph-based methods. Diverse factors can be connected through patients.

Table 1

Notations.

Variable	Interpretation
\mathcal{N}, \mathcal{R}	the set of entities and relations
(h, r, t)	a triple in knowledge graph
\mathbf{h}, \mathbf{t}	both are k dimensional vector embeddings
\mathbf{r}	a d dimensional relation embedding
\mathbf{H}_r	a $k \times d$ dimensional projection matrix
p, m, d	a patient, medicine, disease
$\mathbf{p}, \mathbf{m}, \mathbf{d}$	a k dimensional vector of a patient, a medicine or a disease
$\mathbb{R}^k, \mathbb{R}^d$	k and d dimensional latent space of entities and relations

2. The proposed framework

In this section, we will first describe the notations and formulate medicine recommendation problem, and then present graph embedding models and how to use learned embeddings to recommend safe medicines for patients.

2.1. Problem formulation

Before we focus on the medicine recommendation problem, we first briefly introduce the important notations employed in the remainder of this paper. Table 1 also summarizes them.

The medical knowledge graph describes the medical entities collected from the integrated sources, as well as relationships among these entities. For instance, a triple (*glucocorticoid, adverse interaction, aspirin*) indicates that there is a relationship *adverse interaction* from *glucocorticoid* to *aspirin* in DrugBank. We define the medical knowledge graph as follows.

Definition 1 (Medical knowledge graph). The medical knowledge graph $G = (\mathcal{N}, \mathcal{R})$ is a set of triples in the form (h, r, t) , where \mathcal{N} is a set of entities, \mathcal{R} is a set of relations, $h, t \in \mathcal{N}$ and $r \in \mathcal{R}$.

To capture the co-relationships of patients, diseases, and medicines in EMRs, we define the patient-disease, patient-medicine bipartite graphs as follows.

Definition 2 (Patient-medicine bipartite graph). The patient-medicine bipartite graph is denoted as $G_{pm} = (\mathcal{P} \cup \mathcal{M}, \mathcal{E}_{pm})$, where \mathcal{P} is a set of patients and \mathcal{M} is a set of medicines. \mathcal{E}_{pm} is the set of edges. If a patient p_i takes a medicine m_j , there will be an edge e_{ij} between them, otherwise none. The weight w_{ij} of the edge between patient p_i and medicine m_j is defined as the total times of patient p_i takes the medicine m_j .

Definition 3 (Patient-disease bipartite graph). The patient-disease bipartite graph is denoted as $G_{pd} = (\mathcal{P} \cup \mathcal{D}, \mathcal{E}_{pd})$, where \mathcal{P} is a set

of patients and \mathcal{D} is a set of diseases. \mathcal{E}_{pd} is the set of edges. If a patient p_i is diagnosed with a disease d_j , there will be an edge e_{ij} between them, otherwise none. The weight w_{ij} is set to 1 when the edge e_{ij} exists.

Fig. 2 illustrates a heterogeneous graph by constructing patient-disease, patient-medicine bipartite graphs from MIMIC-III, and linking them to medical knowledge graphs, ICD-9 ontology, and DrugBank. Finally, we formally define the safe medicine recommendation problem as follows.

Problem 1 (Safe medicine recommendation). Given a patient p and his/her diagnoses dataset \mathcal{D}_p , recommending safe medicines for each $d \in \mathcal{D}_p$ is predicting edges from p to medicines dataset \mathcal{M} . The output is a set of medicines \mathcal{M}_p with minimum drug-drug interactions.

2.2. Model description and optimization

In this section, we propose embedding learning approaches to encode the heterogeneous graph in the latent space and its optimization method.

2.2.1. Medical knowledge graph embedding

A medical knowledge graph $G = (\mathcal{N}, \mathcal{R})$ is a multi-relational graph, in which entities \mathcal{N} and relations \mathcal{R} can be different types. For a triple, $(h, r, t) \in G$, we use bold letters $\mathbf{h}, \mathbf{r}, \mathbf{t}$ to denote the corresponding embedding representations of h, r, t . Plenty of graph embedding methods have been proposed to encode a multi-relational graphs into a continuous vector space. Translation-based models [10–12] regard the relation r in each (h, r, t) as a translation from h to t within the low dimensional space, i.e., $\mathbf{h} + \mathbf{r} - \mathbf{t}$, and perform much more effectively and efficiently than conventional models. TransR [12] is a state-of-the-art translation-based embedding approach. It represents entities and relations in distinct vector space bridged by relation-specific matrices to get better graph representations.

Consider the above reason, we set entities embeddings $\mathbf{h}, \mathbf{t} \in \mathbb{R}^k$ and relations embeddings $\mathbf{r} \in \mathbb{R}^d$. And we set a projection matrix $\mathbf{H}_r \in \mathbb{R}^{k \times d}$, which projects entities from entity space to relation space. We define the translations between entities and get the energy function $z(h, r, t)$ as:

$$z(\mathbf{h}, \mathbf{r}, \mathbf{t}) = b - \|\mathbf{h}\mathbf{H}_r + \mathbf{r} - \mathbf{t}\mathbf{H}_r\|_{L1/L2} \quad (1)$$

where b is a bias constant.

Then, the conditional probability of a triple (h, r, t) is defined as follows:

$$P(h|r, t) = \frac{\exp\{z(\mathbf{h}, \mathbf{r}, \mathbf{t})\}}{\sum_{\hat{h} \in \mathcal{N}} \exp\{z(\hat{\mathbf{h}}, \mathbf{r}, \mathbf{t})\}} \quad (2)$$

and $P(t|h, r)$, $P(r|h, t)$ can be defined in the analogous manner. We define the likelihood of observing a triple (h, r, t) as:

$$\begin{aligned} \mathcal{L}(h, r, t) &= \log P(h|r, t) + \log P(t|h, r) \\ &\quad + \log P(r|h, t) \end{aligned} \quad (3)$$

We define an objective function by maximizing the conditional likelihoods of existing triples in G :

$$\mathcal{L}_G = \sum_{(h, r, t) \in G} \mathcal{L}(h, r, t) \quad (4)$$

Based on Eq. (4), the objective functions of medicine and disease knowledge graph $G_m = (\mathcal{N}_m, \mathcal{R}_m)$, $G_d = (\mathcal{N}_d, \mathcal{R}_d)$ can be defined respectively:

$$\mathcal{L}_{G_m} = \sum_{(h_m, r_m, t_m) \in G_m} \mathcal{L}(h_m, r_m, t_m) \quad (5)$$

$$\mathcal{L}_{G_d} = \sum_{(h_d, r_d, t_d) \in G_d} \mathcal{L}(h_d, r_d, t_d) \quad (6)$$

2.2.2. Bipartite graph embedding

Different from the medical knowledge graph, the patient-disease, patient-medicine are bipartite graphs. A bipartite graph has only one single type of relations. For a bipartite graph, LINE [13] model achieves the state-of-the-art performance of encoding the entities into a continuous vector space while preserving co-relations information of the graph. Hence, we follow LINE and set patients, medicines, and diseases embeddings $\mathbf{p}, \mathbf{m}, \mathbf{d} \in \mathbb{R}^k$. We present the process of encoding patient-medicine bipartite graph as follows.

Given a patient-medicine bipartite graph $G_{pm} = (\mathcal{P} \cup \mathcal{M}, \mathcal{E}_{pm})$. We first define the conditional probability of that a patient p_i in set \mathcal{P} takes medicine m_j in set \mathcal{M} as follows:

$$P(m_j|p_i) = \frac{\exp\{z(\mathbf{p}_i, \mathbf{m}_j)\}}{\sum_{\hat{m}_j \in \mathcal{P}_i} \exp\{z(\mathbf{p}_i, \hat{\mathbf{m}}_j)\}} \quad (7)$$

where $z(\mathbf{p}_i, \mathbf{m}_j) = \mathbf{m}_j^T \cdot \mathbf{p}_i$, \mathbf{p}_i is the embedding vector of the patient p_i in \mathcal{P} , and \mathbf{m}_j is the embedding vector of medicine m_j in \mathcal{M} . Eq. (7) defines a conditional distribution $P(\cdot|p_i)$ over all medicines in \mathcal{M} . The empirical distribution $\hat{P}(\cdot|p_i)$ is defined as $\hat{P}(m_j|p_i) = \frac{w_{ij}}{\text{sum}_i}$, where w_{ij} is the weight of the edge e_{ij} and $\text{sum}_i = \sum_j w_{ij}$ is the total times that the patient p_i takes medicines. We maximize the following objective function:

$$\mathcal{L}_{G_{pm}} = - \sum_{p_i \in \mathcal{P}} \lambda_i d(\hat{P}(\cdot|p_i), P(\cdot|p_i)) \quad (8)$$

where $d(\cdot, \cdot)$ is the distance between two distributions. In this paper, we use KL-divergence to compute $d(\cdot, \cdot)$. As sum_i is different from patients, we use $\lambda_i = \text{sum}_i$ in the objective function to represent the personalization of the patient p_i in the graph. After omitting some constants, we have:

$$\mathcal{L}_{G_{pm}} = \sum_{e_{ij} \in \mathcal{E}_{pm}} w_{ij} \log(P(m_j|p_i)) \quad (9)$$

For the patient-disease bipartite graph, we can get the object function $\mathcal{L}_{G_{pd}}$ in the analogous manner:

$$\mathcal{L}_{G_{pd}} = \sum_{e_{ij} \in \mathcal{E}_{pd}} w_{ij} \log(P(d_j|p_i)) \quad (10)$$

2.2.3. Optimization and training

To learn the medical knowledge graph and bipartite graphs embeddings simultaneously, an intuitive approach is to collectively embed the four graphs (patient-medicine bipartite graph, patient-disease bipartite graph, medicine knowledge graph, disease knowledge graph) by maximizing the sum of the four logarithm likelihood objective functions just as follows:

$$\mathcal{L}(X) = \mathcal{L}_{G_m} + \mathcal{L}_{G_d} + \mathcal{L}_{G_{pm}} + \mathcal{L}_{G_{pd}} + \gamma C(X) \quad (11)$$

where X stands for the embeddings $\mathbb{R}_k, \mathbb{R}_d$ of entities and relations in the heterogeneous graph we construct, γ is a hyper-parameter weighting the regularization factor $C(X)$, which is defined as follows:

$$\begin{aligned} C(X) &= \sum_{n_m \in \mathcal{N}_m} [|n_m| - 1]_+ + \sum_{n_d \in \mathcal{N}_d} [|n_d| - 1]_+ \\ &\quad + \sum_{r_m \in \mathcal{R}_m} [|r_m| - 1]_+ + \sum_{r_d \in \mathcal{R}_d} [|r_d| - 1]_+ \\ &\quad + \sum_{p \in \mathcal{P}} [|p| - 1]_+ + \sum_{d \in \mathcal{D}} [|d| - 1]_+ + \sum_{m \in \mathcal{M}} [|m| - 1]_+ \end{aligned} \quad (12)$$

where $[x]_+ = \max(0, x)$ denotes the positive part of x . The regularization factor will normalize the embeddings during learning. And we adopt the asynchronous stochastic gradient algorithm (ASGD) [14] to maximize the transformed objective function.

Optimizing objective functions Eq. (5), Eq. (6), Eq. (9) and Eq. (10) in Eq. (11) are computationally expensive, as calculating them need to sum over the entire set of entities and relations. To address this problem, we use the negative sampling method [15] to transform the objective functions.

For Eq. (5) and Eq. (6), we should transform $\log P(t|h, r)$, $\log P(r|h, t)$, $\log P(h|r, t)$ in Eq. (3). Taking $P(t|h, r)$ as an example, we maximize the following objective function instead of it:

$$\begin{aligned} &\log \sigma(z(\mathbf{h}, \mathbf{r}, \mathbf{t})) \\ &+ \sum_{n=1}^{C_1} E_{\tilde{h}_n \sim z_{neg}(\{(h, r, t)\})} [\sigma(z(\tilde{\mathbf{h}}_n, \mathbf{r}, \mathbf{t}))] \end{aligned} \quad (13)$$

where C_1 is the number of negative examples, $\sigma(x) = 1/(1 + \exp(-x))$ is the sigmoid function. $\{(h, r, t)\}$ is the invalid triple set, and z_{neg} is a function randomly sampling instances from $\{(h, r, t)\}$. When a positive triple $(h, r, t) \in G$ is selected, to maximize Eq. (13), C_1 negative triples are constructed by sampling entities from an uniform distribution over \mathcal{N} and replacing the head of (h, r, t) . The transformed objective of $\log P(r|h, t)$, $\log P(t|h, r)$ are maximized in the same manner, but for $\log P(r|h, t)$, the negative relations are sampled from a uniform distribution over \mathcal{R} to corrupt the positive relation $r \in (h, r, t)$. We iteratively select random mini-batch from the training set to learn embeddings until converge.

For Eq. (9), we also use the negative sampling method to transform it to the following objective function:

$$\begin{aligned} &\log \sigma(z(\mathbf{p}_i, \mathbf{m}_j)) \\ &+ \sum_{n=1}^{C_2} E_{\tilde{m}_n \sim z_{neg}(\tilde{m})} [\log \sigma(z(\mathbf{p}_i, \tilde{\mathbf{m}}_n))] \end{aligned} \quad (14)$$

where $\sigma(x) = 1/(1 + \exp(-x))$ is the sigmoid function, C_2 is the number of negative edges. $z_{neg}(\tilde{m}) \propto \text{sum}_{\tilde{m}}^{3/4}$ according to the empirical setting of [15], $\text{sum}_{\tilde{m}}$ is the total number of times that the medicine \tilde{m} is taken by patients. We can simplify Eq. (10) and maximize it in the same way.

Finally, we can efficiently learn the embeddings of different types of parts in the heterogeneous graph.

2.3. Safe medicine recommendation process

In this section, we present how to recommend safe medicines based on the learned embeddings and diagnoses of a given patient. For an existing patient p , we use the learned embedding \mathbf{p} to predict new medicine recommendations. For a new given patient p , we first use diseases embeddings of p 's diagnoses to represent \mathbf{p} , and then recommend safe medicines for p , as shown in Fig. 2.

2.3.1. New patient model

We aim to present a new patient p by his/her diagnoses embeddings. We should consider the time sequence of diseases that a patient is diagnosed, especially for the patient with multiple diseases. Assume a patient p in the hospital or on medication is associated with n ranked diseases according to their timestamps in an increasing order. Then, the patient embedding \mathbf{p} can be encoded as follows:

$$\mathbf{p} = \sum_{t=1}^n \exp^{-t} \cdot \mathbf{d}_t \quad (15)$$

where \mathbf{d}_t is the t -th embedding of disease d_t .

2.3.2. Medicine recommendation

Given a query patient p with the query disease d , i.e., $q = (p, d)$, we first project disease d and patient p into their latent space, and then select top- k safe medicines.¹ More precisely, given a query $q = (p, d)$, for each medicine m which could be useful for p , we compute its ranking score as in Eq. (16), and then select the medicine m with the top- k highest ranking scores as the recommendation.

$$S(q, m_n) = \mathbf{p}^T \cdot \mathbf{m}_n - \sum_{o=1}^{n-1} \|\mathbf{m}_n + \mathbf{r}_{\text{interaction}} - \mathbf{m}_o\|_{L1/L2} \quad (16)$$

where \mathbf{p} is the representation of patient p and m_n is the n -th medicines to be considered from medicines \mathcal{M} based on the already selected medicine m_1, \dots, m_{n-1} .

3. Experiments and evaluation

We attempt to demonstrate the effectiveness of our recommendation method in this section, which is referred to as SMR in this paper. In particular, we expect to answer “**how well does our method compare with the competing techniques?**” in Section 3.2. The results show that our recommendation method significantly outperforms the three baselines. The detailed experimental settings of our evaluations are described in Section 3.1.

3.1. Experimental settings

3.1.1. Data sets

Our experiments are performed on the real EMRs datasets, MIMIC-III [16], and two medical knowledge graphs, ICD-9 ontology [9] and DrugBank [8]. These real datasets are publicly available in different forms.

¹ In MIMIC-III, patients in ICUs are sicker and usually need more medicines for a diagnosis. We set $k = 3$ in this paper.

Table 2
Entities and relations in the heterogeneous graph.

Entities	Relations	
#Disease	6,985	#Medicine-related
#Medicine	8,054	#Disease-related
#Patient	46,520	#Patient-Disease
#Medicine-related	305,642	#Patient-Medicine

Table 3
Experiments on medicine group 1.

	Prediction accuracy	DDIs rate
Rule-based	0.3068	32.01%
K-Most frequent	0.3473	14.08%
LEAP	0.5582	1.66%
SMR	0.6113	0.17%

- MIMIC-III (Medical Information Mart for Intensive Care III) collected bedside monitor trends, electronic medical notes, laboratory test results, and waveforms from the ICUs (Intensive Care Units) of Beth Israel Deaconess Medical Center between 2001 and 2012. It contains distinct 46,520 patients, 650,987 diagnoses and 1,517,702 prescription records that associated with 6,985 distinct diseases and 4,525 medicines.
- ICD-9 ontology² (International classification of diseases-version 9) contains 13,000 international standard codes of diagnoses and the relationships between them.
- DrugBank is a bioinformatics/cheminformatics resource which consists of medicine related entities. The medical knowledge graph version³ contains 8,054 medicines, 4,038 other related entities (e.g., protein or drug targets) and 21 relationships.

Heterogeneous graph construction We connect MIMIC-III, ICD-9 ontology, and DrugBank (medicine group 1) by constructing the patient-medicine bipartite graph and the patient-disease bipartite graph.

For the patient-disease bipartite graph, MIMIC-III provides ICD-9 codes for diagnoses, which implicitly the diagnoses of MIMIC-III can be linked to ICD-9 ontology by string matching. For the patient-medicine bipartite graph, the prescriptions in MIMIC-III consist of the drug information, e.g., the names, the duration, and the dosage. However, various names to a single type of medicine in MIMIC-III exist due to some noisy words (20%, 50 ml, glass bottle, etc.), which becomes an obstacle to link medicine names to DrugBank when directly applying to the string matching method. We use an entity linking method [17] instead to address this problem. Table 2 shows the statistic of the heterogeneous graph we construct. The heterogeneous graph will be used to learn low-dimension representations of entities and relations by the SMR framework. Afterward, we categorize the medicines in the heterogeneous graph into two groups: 1). The first group consists of all 4,525 medicines that are recorded in EMRs, and will be used as inputs of the baseline methods. 2). The second group contains 3,529 medicines that haven't been observed in EMRs, and will be used as test data for cold start recommendation.

3.1.2. Baselines

We compare our SMR with the following state of the art methods:

- Rule-based method [2] recommends medicines based on mappings from existing medicine categories to diseases in the

² <http://bioportal.bioontology.org/ontologies/ICD9CM>.

³ <http://wif05-03.informatik.uni-mannheim.de/drugbank/>.

1 **Table 4**
2 Experiments on medicine group 2.
3

	Prediction accuracy	DDIs rate
Rule-based	0.2736	27.01%
K-Most frequent	NA	NA
LEAP	NA	NA
SMR	0.5214	2.01%

9 MEDI database [18]. For each disease, a drug is assigned to
10 the patient according to the mappings.

- 11 • K-Most frequent method is a basic baseline which retrieves
12 the top K medicines that most frequently co-occur with each
13 disease as their recommendation. We set $K = 3$ in this paper.
- 14 • LEAP method [3] uses a Multi-Instance Multi-Label learning
15 framework to train a predictive model taking disease conditions
16 as input features and yielding multiple medicine labels
17 as recommendations.

18 3.2. Evaluation methods

19 To guarantee medicine recommendations generated by SMR
20 work effectively, we evaluate four indices, the prediction accuracy,
21 the ability to avoid adverse drug-drug interactions, the experienced
22 clinical doctor assessments, and the capacity to process cold start
23 problem. In all experiments, the ratio of training to validation to
24 test sets is 0.7:0.1:0.2. The hyper-parameters was adjusted by a vali-
25 dation set.

26 3.2.1. Prediction accuracy and DDIs rate

27 We utilize Jaccard Coefficient to compare the similarity of the
28 prescriptions generated by SMR and the corresponding prescrip-
29 tions written by doctors. Given the recommendation medicines set
30 M_i generated by SMR for a patient p_i , \check{M}_i is the medicines set pre-
31 scribed by doctors in the data. The mean of Jaccard coefficient is
32 defined as follows,

$$33 \quad Jaccard = \frac{1}{K} \sum_{i=1}^K \frac{|M_i \cap \check{M}_i|}{|M_i \cup \check{M}_i|} \quad (17)$$

34 where K is the number of samples in a test set. Table 3 shows
35 the accuracy of the baselines and SMR on medicine group 1, the
36 rule-based method performs the worst because it is the only one
37 provides a general recommendation for a specific diagnosis and
38 it is not able to endow personalized recommendations, especially
39 for the patients with multiple diseases. The frequency of each
40 medicine-disease pair remains high in ICUs. Hence, recom-
41 mendations based on frequency, K-Most frequent method, also work
42 deficiently. Our method SMR outperforms LEAP by 1.49% because
43 more accurate medical facts are involved in medical knowledge
44 graphs rather than the prescription information in EMRs.

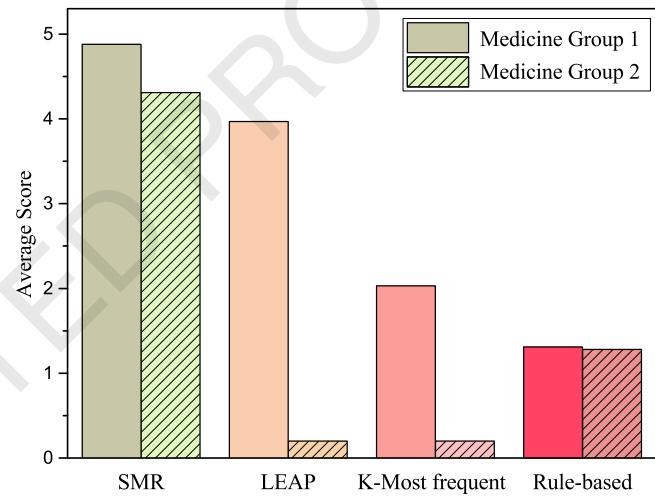
45 We extract all adverse drug-drug interactions (DDIs) from Drug-
46 Bank to evaluate whether medicine recommendations embrace un-
47 safe DDIs. Table 3 shows the percentages of different medicine
48 recommendations consisting of adverse DDIs. The result indicates
49 that SMR can recommend most harmless medicines for patients
50 as its drug interaction rate is the lowest. The rule-based method
51 and K-Most frequent method select medicines by a greedy strat-
52 egy only regardless of specific adverse DDIs. For the rarely used
53 medicines and unknown DDIs in EMRs, SMR is more reliable than
54 LEAP. The reason is that SMR can predict each patient-medicine
55 link and compute potential hidden DDIs by the learned embed-
56 dings of medical knowledge graphs.

57 3.2.2. Cold start

58 We evaluate the ability of baselines and SMR in addressing
59 cold start medicine recommendations on the medicine group 2.

60 **Table 5**
61 Examples of medicine recommendations generated by Baselines and SMR.

Diagnosis	Methods	Medicine recommendations
Sepsis	Rule-based	Teicoplanin, Metoprolol
Acute respiratory failure	K-Most frequent	Vancomycin, Furosemide, Metoprolol, Insulin
Hypertension	LEAP	Vancomycin, Furosemide, Metoprolol Tartrate
	SMR	Vancomycin, Furosemide, Amlodipine, Norepinephrine, Acetaminophen
Type 2 diabetes	Rule-based	Gliclazide, Phenylbutazone, Sulfasalazine
Rheumatoid arthritis	K-Most frequent	Furosemide, Tolbutamide, Phenylbutazone, Metoprolol, Insulin, Acetaminophen
Hypertension	Hyperlipidemia	Metformin, Amethopterin, Amiloride/HCTZ, Fenofibrate
	LEAP	Metformin, Insulin, Acetaminophen, Nifedipine, Fenofibrate
	SMR	



62 **Fig. 3.** Clinical assessment.

63 The results are reported in Table 4. K-Most frequent method and
64 LEAP are not applicable (NA) on recommending new medicines in
65 the cold start scenario. Since our SMR process can present new
66 medicines by the learned vector representations of used medicines,
67 the potential patient-medicine links between cold start medicines
68 and patients will be captured correspondingly. In other words,
69 SMR can leverage not only the patient-medicine links in EMRs but
70 also the medical knowledge graphs when recommending cold start
71 medicines.

72 3.2.3. Clinical assessment

73 We invited three experienced clinical experts to evaluate the
74 effectiveness of the medicine recommendations by scoring on a
75 6-point scale: 5 corresponding to completely cover all diagnoses
76 without DDIs; 4 to partially (at least 50%) diagnoses include without
77 DDIs; 3 to completely cover all diagnoses with DDIs; 2 to less
78 than 50% diagnoses without DDIs; 1 to partially (at least 50%) diag-
79 noses covered with DDIs; 0 to less than 50% diagnoses with DDIs.
80 The average score of three experts is used as the final clinical as-
81 sessment score for each recommendation, as shown in Fig. 3.

82 3.2.4. Case study

83 In Table 5 we illustrate two events of medicine recommendations
84 on medicine group 1 for patients associated with multiple
85 types of diseases. SMR is qualified to succeed in all these two

cases when comparing it against other baselines. For the first patient, SMR recommended a set of medicines with 100% coverage, with Vancomycin for Sepsis, Norepinephrine, Acetaminophen for respiratory failure, Furosemide and Amlodipine for Hypertension. In contrast, other baselines are not capable of make an adequate consideration. The rule-based method adopted Teicoplanin, targeting Sepsis only and not appropriate. The K-Most frequent method and LEAP only selected Vancomycin for Sepsis and other medicines for Hypertension. For the second patient, SMR recommends more suitable medicines than LEAP and Rule-based method, i.e., Metformin and Insulin for Type 2 diabetes, Acetaminophen to release Rheumatoid arthritis, Nifedipine for Hypertension, and Fenofibrate for Hyperlipidemia. There is an adverse DDI among the medicines recommended by the K-Most frequent method. Tolbutamide and Phenylbutazone can lead to harmful, potentially fatal effects when taken together. This case also indicates SMR can avoid the adverse DDIs when recommending medicines.

4. Related work

In this section, we discuss related work, including medicine recommendation, medical knowledge graphs and embeddings.

4.1. Medicine recommendation

As introduced in Section 1, two types of methods, rule-based protocols [1,19,2], and supervised-learning-based methods [3,20], are currently utilizing EMRs to recommending medicines. Ideally, medicine recommendation systems aim to tailor treatment to the individual characteristics of each patient [21]. Hence, medicine recommendation has also received attention recently in genetics/genomics research fields. There are already existing medicine recommendation systems [22,23] by leveraging genetics/genomics information of patients in current practice, such information is not yet widely available in everyday clinical practice, and is insufficient since it only addresses one of many factors affecting response to medication. A previous pre-print version of our work [24] mainly introduced our idea and recommendation model, while in this paper we provide a more in-depth analysis of the experiment and a summary of the relevant work.

4.1.1. Medical knowledge graphs and embeddings

Recent evaluation efforts on knowledge graphs have focused on automatic knowledge base population and completion. Some knowledge graphs have also been constructed from huge volumes of medical databases over the last years, such as [25], Bio2RDF [26], and Chem2Bio2RDF [27]. Medical knowledge graphs contain an abundance of basic medical facts of medicines and diseases and provide a pathway for medical discovery and applications, such as effective medicine recommendations. Unfortunately, such medical knowledge graphs suffer from serious data incomplete problem, which impedes its application in the field of clinical medicine. Recently, Celebi et al. [28] proposed a knowledge graph embedding approach for drug-drug interaction prediction in realistic scenario. MedGraph [29] is a new EMR embedding framework which introduces a graph-based data structure to capture both structural visit-code co-location information structurally and temporal visit sequencing information. However, they only focused on the real-world biomedical embeddings learning and do not directly harness the medicine recommendation task, thus the specific clinical medicine information is not sufficiently tailored. In contrast, our model jointly embeds diseases, medicines, patients, and uses the learned embeddings to decompose the medicine recommendation into a linked prediction process while considering the patient's diagnoses and adverse drug reactions.

5. Conclusion and future work

In this paper, we propose a novel framework SMR to recommend safe medicines for patients, especially for the patients with multiple diseases. SMR first constructs a high-quality heterogeneous graph by bridging EMRs (MIMIC-III) and medical knowledge graphs (ICD-9 ontology and DrugBank). Then, SMR jointly embeds diseases, medicines, patients, and their corresponding relations into a shared lower dimensional space. Finally, SMR uses the embeddings to decompose the medicine recommendation into a linked prediction process considering the clinical diagnoses and adverse DDI reactions. Extensive experiments on real world datasets are conducted and demonstrate the effectiveness of SMR. In future work, we will improve the linking accuracy by considering more information of patients, such as the clinical outcomes and demographics.

6. Acknowledgment

This work was supported by National Science Foundation of China with Grant Nos. 61906037 and U1736204; the Fundamental Research Funds for the Central Universities with Grant Nos. 4009009106 and 22120200184; the CCF-Baidu Open Fund.

Declaration of competing interest

The authors declare the following financial interests/personal relationships which may be considered as potential competing interests.

References

- [1] Z. Chen, K. Marple, E. Salazar, G. Gupta, L. Tamil, A physician advisory system for chronic heart failure management based on knowledge patterns, *Theory Pract. Log. Program.* 16 (5–6) (2016) 604–618.
- [2] D. Almirall, S.N. Compton, M. Gunlicks-Stoessel, N. Duan, S.A. Murphy, Designing a pilot sequential multiple assignment randomized trial for developing an adaptive treatment strategy, *Stat. Med.* 31 (17) (2012) 1887–1902.
- [3] Y. Zhang, R. Chen, J. Tang, W.F. Stewart, J. Sun Leap, Learning to prescribe effective and safe treatment combinations for multimorbidity, in: Proceedings of the 23rd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining, ACM, 2017, pp. 1315–1324.
- [4] M. Panagioti, J. Stokes, A. Esmail, P. Coventry, S. Cheraghi-Sohi, R. Alam, P. Bower, Multimorbidity and patient safety incidents in primary care: a systematic review and meta-analysis, *PLoS ONE* 10 (8) (2015) e0135947.
- [5] D.N. Juurlink, M. Mamdani, A. Kopp, A. Laupacis, D.A. Redelmeier, Drug-drug interactions among elderly patients hospitalized for drug toxicity, *J. Am. Med. Assoc.* 289 (13) (2003) 1652–1658.
- [6] I.R. Edwards, J.K. Aronson, Adverse drug reactions: definitions, diagnosis, and management, *Lancet* 356 (9237) (2000) 1255–1259.
- [7] L.L. Leape, D.W. Bates, D.J. Cullen, J. Cooper, H.J. Demonaco, T. Gallivan, R. Hallisey, J. Ives, N. Laird, G. Laffel, et al., Systems analysis of adverse drug events, *J. Am. Med. Assoc.* 274 (1) (1995) 35–43.
- [8] V. Law, C. Knox, Y. Djoumbou, T. Jewison, A.C. Guo, Y. Liu, A. Maciejewski, D. Arndt, M. Wilson, V. Neveu, et al., Drugbank 4.0: shedding new light on drug metabolism, *Nucleic Acids Res.* 42 (D1) (2014) D1091–D1097.
- [9] L.M. Schriml, C. Arze, S. Nadendla, Y.-W.W. Chang, M. Mazaitis, V. Felix, G. Feng, W.A. Kibbe, Disease ontology: a backbone for disease semantic integration, *Nucleic Acids Res.* 40 (D1) (2011) D940–D946.
- [10] A. Bordes, N. Usunier, A. Garcia-Duran, J. Weston, O. Yakhnenko, Translating embeddings for modeling multi-relational data, in: Advances in Neural Information Processing Systems, 2013, pp. 2787–2795.
- [11] Z. Wang, J. Zhang, J. Feng, Z. Chen, Knowledge graph embedding by translating on hyperplanes, in: Proceedings of the 28th AAAI Conference on Artificial Intelligence, 2014, pp. 1112–1119.
- [12] Y. Lin, Z. Liu, M. Sun, Y. Liu, X. Zhu, Learning entity and relation embeddings for knowledge graph completion, in: Proceedings of the 29th AAAI Conference on Artificial Intelligence, 2015, pp. 2181–2187.
- [13] J. Tang, M. Qu, M. Wang, M. Zhang, J. Yan, Q. Mei Line, Large-scale information network embedding, in: Proceedings of the 24th International Conference on World Wide Web, International World Wide Web Conferences Steering Committee, 2015, pp. 1067–1077.

- 1 [14] B. Recht, C. Re, S. Wright, F. Niu, Hogwild: a lock-free approach to parallelizing stochastic gradient descent, in: Advances in Neural Information Processing Systems, 2011, pp. 693–701.
- 2 [15] T. Mikolov, I. Sutskever, K. Chen, G.S. Corrado, J. Dean, Distributed representations of words and phrases and their compositionality, in: Advances in Neural Information Processing Systems, 2013, pp. 3111–3119.
- 3 [16] A.E. Johnson, T.J. Pollard, L. Shen, L.-w.H. Lehman, M. Feng, M. Ghassemi, B. Moody, P. Szolovits, L.A. Celi, R.G. Mark, Mimic-iii, a freely accessible critical care database, *Sci. Data* 3 (2016).
- 4 [17] M. Wang, J. Zhang, J. Liu, W. Hu, S. Wang, X. Li, W. Liu, Pdd graph: bridging electronic medical records and biomedical knowledge graphs via entity linking, in: International Semantic Web Conference, Springer, 2017.
- 5 [18] W.-Q. Wei, R.M. Cronin, H. Xu, T.A. Lasko, L. Bastarache, J.C. Denny, Development and evaluation of an ensemble resource linking medications to their indications, *J. Am. Med. Inform. Assoc.* 20 (5) (2013) 954–961.
- 6 [19] M. Gunlicks-Stoessel, L. Mufson, A. Westervelt, D. Almirall, A pilot smart for developing an adaptive treatment strategy for adolescent depression, *J. Clin. Child Adolesc. Psychol.* 45 (4) (2016) 480–494.
- 7 [20] P. Zhang, F. Wang, J. Hu, R. Sorrentino, Towards personalized medicine: leveraging patient similarity and drug similarity analytics, *Proc. AMIA Summits Transl. Sci.* 2014 (2014) 132.
- 8 [21] G.H. Fernald, E. Capriotti, R. Daneshjou, K.J. Karczewski, R.B. Altman, Bioinformatics challenges for personalized medicine, *Bioinformatics* 27 (13) (2011) 1741–1748.
- 9 [22] M. Rosen-Zvi, A. Altmann, M. Prosperi, E. Aharoni, H. Neuvirth, A. Sönnnerborg, E. Schüller, D. Struck, Y. Peres, F. Incardona, et al., Selecting anti-HIV thera-
- 10 pies based on a variety of genomic and clinical factors, *Bioinformatics* 24 (13) (2008) i399–i406.
- 11 [23] C.C. Bennett, K. Hauser, Artificial intelligence framework for simulating clinical decision-making: a Markov decision process approach, *Artif. Intell. Med.* 57 (1) (2013) 9–19.
- 12 [24] M. Wang, M. Liu, J. Liu, S. Wang, G. Long, B. Qian, Safe medicine recommendation via medical knowledge graph embedding, arXiv preprint, arXiv: 1710.05980, 2017.
- 13 [25] P. Ernst, A. Siu, G. Weikum, Knowlife: a versatile approach for constructing a large knowledge graph for biomedical sciences, *BMC Bioinform.* 16 (1) (2015) 157.
- 14 [26] M. Dumontier, A. Callahan, J. Cruz-Toledo, P. Ansell, V. Emonet, F. Belleau, A. Droit, Bio2rdf release 3: a larger connected network of linked data for the life sciences, in: Proceedings of the 2014 International Conference on Posters & Demonstrations Track-Volume 1272, CEUR-WS.org, 2014, pp. 401–404.
- 15 [27] B. Chen, X. Dong, D. Jiao, H. Wang, Q. Zhu, Y. Ding, D.J. Wild, Chem2bio2rdf: a semantic framework for linking and data mining chemogenomic and systems chemical biology data, *BMC Bioinform.* 11 (1) (2010) 255.
- 16 [28] R. Celebi, H. Uyar, E. Yasar, O. Gumus, O. Dikenelli, M. Dumontier, Evaluation of knowledge graph embedding approaches for drug-drug interaction prediction in realistic settings, *BMC Bioinform.* 20 (1) (2019) 1–14.
- 17 [29] B. Hettige, Y.-F. Li, W. Wang, S. Le, W. Buntine, Medgraph: structural and temporal representation learning of electronic medical records, in: Proceedings of the 24th European Conference on Artificial Intelligence, 2020.

UNCORRECTED PROOF